
Lambda Expressions

The biggest language change in Java 8 is the introduction of lambda expressions—a compact way of passing around behavior. They are also a pretty fundamental building block that the rest of this book depends upon, so let's get into what they're all about.

Your First Lambda Expression

Swing is a platform-agnostic Java library for writing graphical user interfaces (GUIs). It has a fairly common idiom in which, in order to find out what your user did, you register an *event listener*. The event listener can then perform some action in response to the user input (see [Example 2-1](#)).

Example 2-1. Using an anonymous inner class to associate behavior with a button click

```
button.addActionListener(new ActionListener() {  
    public void actionPerformed(ActionEvent event) {  
        System.out.println("button clicked");  
    }  
});
```

In this example, we're creating a new object that provides an implementation of the `ActionListener` class. This interface has a single method, `actionPerformed`, which is called by the button instance when a user actually clicks the on-screen button. The anonymous inner class provides the implementation of this method. In [Example 2-1](#), all it does is print out a message to say that the button has been clicked.



This is actually an example of using *code as data*—we're giving the button an object that represents an action.

Anonymous inner classes were designed to make it easier for Java programmers to pass around code as data. Unfortunately, they don't make it easy enough. There are still four lines of boilerplate code required in order to call the single line of important logic. Look how much gray we get if we color out the boilerplate:

```
button.addActionListener(new ActionListener() {  
    public void actionPerformed(ActionEvent event) {  
        System.out.println("button clicked");  
    }  
});
```

Boilerplate isn't the only issue, though: this code is fairly hard to read because it obscures the programmer's intent. We don't want to pass in an object; what we really want to do is pass in some behavior. In Java 8, we would write this code example as a lambda expression, as shown in [Example 2-2](#).

Example 2-2. Using a lambda expression to associate behavior with a button click

```
button.addActionListener(event -> System.out.println("button clicked"));
```

Instead of passing in an object that implements an interface, we're passing in a block of code—a function without a name. `event` is the name of a parameter, the same parameter as in the anonymous inner class example. `->` separates the parameter from the body of the lambda expression, which is just some code that is run when a user clicks our button.

Another difference between this example and the anonymous inner class is how we declare the variable `event`. Previously, we needed to explicitly provide its type—`ActionEvent event`. In this example, we haven't provided the type at all, yet this example still compiles. What is happening under the hood is that `javac` is inferring the type of the variable `event` from its context—here, from the signature of `addActionListener`. What this means is that you don't need to explicitly write out the type when it's obvious. We'll cover this inference in more detail soon, but first let's take a look at the different ways we can write lambda expressions.



Although lambda method parameters require less boilerplate code than was needed previously, they are still statically typed. For the sake of readability and familiarity, you have the option to include the type declarations, and sometimes the compiler just can't work it out!

How to Spot a Lambda in a Haystack

There are a number of variations of the basic format for writing lambda expressions, which are listed in [Example 2-3](#).

Example 2-3. Some different ways of writing lambda expressions

```
Runnable noArguments = () -> System.out.println("Hello World"); ❶

ActionListener oneArgument = event -> System.out.println("button clicked"); ❷

Runnable multiStatement = () -> { ❸
    System.out.print("Hello");
    System.out.println(" World");
};

BinaryOperator<Long> add = (x, y) -> x + y; ❹

BinaryOperator<Long> addExplicit = (Long x, Long y) -> x + y; ❺
```

❶ shows how it's possible to have a lambda expression with no arguments at all. You can use an empty pair of parentheses, `()`, to signify that there are no arguments. This is a lambda expression implementing `Runnable`, whose only method, `run`, takes no arguments and is a void return type.

❷ we have only one argument to the lambda expression, which lets us leave out the parentheses around the arguments. This is actually the same form that we used in [Example 2-2](#).

Instead of the body of the lambda expression being just an expression, in ❸ it's a full block of code, bookended by curly braces `{ }`. These code blocks follow the usual rules that you would expect from a method. For example, you can return or throw exceptions to exit them. It's also possible to use braces with a single-line lambda, for example to clarify where it begins and ends.

Lambda expressions can also be used to represent methods that take more than one argument, as in ❹. At this juncture, it's worth reflecting on how to *read* this lambda expression. This line of code doesn't add up two numbers; it creates a function that adds together two numbers. The variable called `add` that's a `BinaryOperator<Long>` isn't the result of adding up two numbers; it is code that adds together two numbers.

So far, all the types for lambda expression parameters have been inferred for us by the compiler. This is great, but it's sometimes good to have the option of explicitly writing the type, and when you do that you need to surround the arguments to the lambda expression with parentheses. The parentheses are also necessary if you've got multiple arguments. This approach is demonstrated in ❺.



The *target type* of a lambda expression is the type of the context in which the lambda expression appears—for example, a local variable that it's assigned to or a method parameter that it gets passed into.

What is implicit in all these examples is that a lambda expression's type is context dependent. It gets inferred by the compiler. This target typing isn't entirely new, either. As shown in [Example 2-4](#), the types of array initializers in Java have always been inferred from their contexts. Another familiar example is `null`. You can know what the type of `null` is only once you actually assign it to something.

Example 2-4. The righthand side doesn't specify its type; it is inferred from the context

```
final String[] array = { "hello", "world" };
```

Using Values

When you've used anonymous inner classes in the past, you've probably encountered a situation in which you wanted to use a variable from the surrounding method. In order to do so, you had to make the variable `final`, as demonstrated in [Example 2-5](#). Making a variable `final` means that you can't reassign to that variable. It also means that whenever you're using a `final` variable, you know you're using a specific value that has been assigned to the variable.

Example 2-5. A final local variable being captured by an anonymous inner class

```
final String name = getUserName();
button.addActionListener(new ActionListener() {
    public void actionPerformed(ActionEvent event) {
        System.out.println("hi " + name);
    }
});
```

This restriction is relaxed a bit in Java 8. It's possible to refer to variables that aren't `final`; however, they still have to be *effectively final*. Although you haven't declared the variable(s) as `final`, you still cannot use them as nonfinal variable(s) if they are to be used in lambda expressions. If you do use them as nonfinal variables, then the compiler will show an error.

The implication of being effectively `final` is that you can assign to the variable only once. Another way to understand this distinction is that lambda expressions capture *values*, not variables. In [Example 2-6](#), `name` is an effectively `final` variable.

Example 2-6. An effectively final variable being captured by an anonymous inner class

```
String name = getUserName();
button.addActionListener(event -> System.out.println("hi " + name));
```

I often find it easier to read code like this when the `final` is left out, because it can be just line noise. Of course, there are situations where it can be easier to understand code with an explicit `final`. Whether to use the effectively `final` feature comes down to personal choice.

If you assign to the variable multiple times and then try to use it in a lambda expression, you'll get a compile error. For example, [Example 2-7](#) will fail to compile with the error message: local variables referenced from a lambda expression must be final or effectively final.

Example 2-7. Fails to compile due to the use of a not effectively final variable

```
String name = getUserName();
name = formatUserName(name);
button.addActionListener(event -> System.out.println("hi " + name));
```

This behavior also helps explain one of the reasons some people refer to lambda expressions as “closures.” The variables that aren’t assigned to are *closed* over the surrounding state in order to bind them to a value. Among the chattering classes of the programming language world, there has been much debate over whether Java really has closures, because you can refer to only effectively final variables. To paraphrase Shakespeare: *A closure by any other name will function all the same*. In an effort to avoid such pointless debate, I’ll be referring to them as “lambda expressions” throughout this book. Regardless of what we call them, I’ve already mentioned that lambda expressions are statically typed, so let’s investigate the types of lambda expressions themselves: these types are called *functional interfaces*.

Functional Interfaces



A functional interface is an interface with a single abstract method that is used as the type of a lambda expression.

In Java, all method parameters have types; if we were passing 3 as an argument to a method, the parameter would be an `int`. So what’s the type of a lambda expression?

There is a really old idiom of using an interface with a single method to represent a method and reusing it. It’s something we’re all familiar with from programming in Swing, and it is exactly what was going on in [Example 2-2](#). There’s no need for new magic to be employed here. The exact same idiom is used for lambda expressions, and we call this kind of interface a *functional interface*. [Example 2-8](#) shows the functional interface from the previous example.

Example 2-8. The ActionListener interface: from an ActionEvent to nothing

```
public interface ActionListener extends EventListener {
    public void actionPerformed(ActionEvent event);
}
```

`ActionListener` has only one abstract method, `actionPerformed`, and we use it to *represent* an action that takes one argument and produces no result. Remember, because `actionPerformed` is defined in an interface, it doesn't actually need the `abstract` keyword in order to be abstract. It also has a parent interface, `EventListener`, with no methods at all.

So it's a functional interface. It doesn't matter what the single method on the interface is called—it'll get matched up to your lambda expression as long as it has a compatible method signature. Functional interfaces also let us give a useful name to the type of the parameter—something that can help us understand what it's used for and aid readability.

The functional interface here takes one `ActionEvent` parameter and doesn't return anything (`void`), but functional interfaces can come in many kinds. For example, they may take two parameters and return a value. They can also use generics; it just depends upon what you want to use them for.

From now on, I'll use diagrams to represent the different kinds of functional interfaces you're encountering. The arrows going into the function represent arguments, and if there's an arrow coming out, it represents the return type. For example, an `ActionListener` would look like [Figure 2-1](#).




Figure 2-1. The `ActionListener` interface showing an `ActionEvent` going in and nothing (`void`) coming out

Over time you'll encounter many functional interfaces, but there is a core group in the Java Development Kit (JDK) that you will see time and time again. I've listed some of the most important functional interfaces in [Table 2-1](#).

Table 2-1. Important functional interfaces in Java

Interface name	Arguments	Returns	Example
<code>Predicate<T></code>	<code>T</code>	<code>boolean</code>	Has this album been released yet?
<code>Consumer<T></code>	<code>T</code>	<code>void</code>	Printing out a value
<code>Function<T,R></code>	<code>T</code>	<code>R</code>	Get the name from an <code>Artist</code> object
<code>Supplier<T></code>	None	<code>T</code>	A factory method
<code>UnaryOperator<T></code>	<code>T</code>	<code>T</code>	Logical not (!)
<code>BinaryOperator<T></code>	<code>(T, T)</code>	<code>T</code>	Multiplying two numbers (*)

I've talked about what types functional interfaces take and mentioned that `javac` can automatically infer the types of parameters and that you can manually provide them, but when do you know whether to provide them? Let's look a bit more at the details of type inference.

Type Inference

There are certain circumstances in which you need to manually provide type hints, and my advice is to do what you and your team find easiest to read. Sometimes leaving out the types removes line noise and makes it easier to see what is going on. Sometimes leaving them in can make it clearer what is going on. I've found that at first they can sometimes be helpful, but over time you'll switch to adding them in only when they are actually needed. You can figure out whether they are needed from a few simple rules that I'll introduce in this chapter.

The type inference used in lambdas is actually an extension of the target type inference introduced in Java 7. You might be familiar with Java 7 allowing you to use a diamond operator that asks `javac` to *infer* the generic arguments for you. You can see this in [Example 2-9](#).

Example 2-9. Diamond inference for variables

```
Map<String, Integer> oldWordCounts = new HashMap<String, Integer>(); ❶  
Map<String, Integer> diamondWordCounts = new HashMap<>(); ❷
```

For the variable `oldWordCounts` ❶ we have explicitly added the generic types, but `diamondWordCounts` ❷ uses the diamond operator. The generic types aren't written out—the compiler just figures out what you want to do by itself. Magic!

It's not really magic, of course. Here, the generic types to `HashMap` can be inferred from the type of `diamondWordCounts` ❷. You still need to provide generic types on the variable that is being assigned to, though.

If you're passing the constructor straight into a method, it's also possible to infer the generic types from that method. In [Example 2-10](#), we pass a `HashMap` as an argument that already has the generic types on it.

Example 2-10. Diamond inference for methods

```
useHashMap(new HashMap<>());  
  
...  
  
private void useHashMap(Map<String, String> values);
```

In the same way that Java 7 allowed you to leave out the generic types for a constructor, Java 8 allows you to leave out the types for whole parameters of lambda expressions.

Again, it's not magic: `javac` looks for information close to your lambda expression and uses this information to figure out what the correct type should be. It's still type checked and provides all the safety that you're used to, but you don't have to state the types explicitly. This is what we mean by *type inference*.



It's also worth noting that in Java 8 the type inference has been improved. The earlier example of passing `new HashMap<>()` into a `use HashMap` method actually wouldn't have compiled in Java 7, even though the compiler had all the information it needed to figure things out.

Let's go into a little more detail on this point with some examples.

In both of these cases we're assigning the variables to a functional interface, so it's easier to see what's going on. The first example ([Example 2-11](#)) is a lambda that tells you whether an `Integer` is greater than 5. This is actually a `Predicate`—a functional interface that checks whether something is true or false.

Example 2-11. Type inference

```
Predicate<Integer> atLeast5 = x -> x > 5;
```

A `Predicate` is also a lambda expression that returns a value, unlike the previous `ActionListener` examples. In this case we've used an expression, `x > 5`, as the body of the lambda expression. When that happens, the return value of the lambda expression is the value its body evaluates to.

You can see from [Example 2-12](#) that `Predicate` has a single generic type; here we've used an `Integer`. The only argument of the lambda expression implementing `Predicate` is therefore inferred as an `Integer`. `javac` can also check whether the return value is a `boolean`, as that is the return type of the `Predicate` method (see [Figure 2-2](#)).

Example 2-12. The predicate interface in code, generating a boolean from an Object

```
public interface Predicate<T> {  
    boolean test(T t);  
}
```




Figure 2-2. The Predicate interface diagram, generating a boolean from an Object

Let's look at another, slightly more complex functional interface example: the `BinaryOperator` interface, which is shown in [Example 2-13](#). This interface takes two arguments and returns a value, all of which are the same type. In the code example we've used, this type is `Long`.

Example 2-13. A more complex type inference example

```
BinaryOperator<Long> addLongs = (x, y) -> x + y;
```

The inference is smart, but if it doesn't have enough information, it won't be able to make the right decision. In these cases, instead of making a wild guess it'll just stop what it's doing and ask for help in the form of a compile error. For example, if we remove some of the type information from the previous example, we get the code in [Example 2-14](#).

Example 2-14. Code doesn't compile due to missing generics

```
BinaryOperator add = (x, y) -> x + y;
```

This code results in the following error message:

```
Operator '&#x002B;' cannot be applied to java.lang.Object, java.lang.Object.
```

That looks messy: what is going on here? Remember that `BinaryOperator` was a functional interface that had a generic argument. The argument is used as the type of both arguments, `x` and `y`, and also for its return type. In our code example, we didn't give any generics to our `add` variable. It's the very definition of a *raw* type. Consequently, our compiler thinks that its arguments and return values are all instances of `java.lang.Object`.

We will return to the topic of type inference and its interaction with method overloading in [“Overload Resolution” on page 45](#), but there's no need to understand more detail until then.

Key Points

- A lambda expression is a method without a name that is used to pass around behavior as if it were data.
- Lambda expressions look like this: `BinaryOperator<Integer> add = (x, y) -> x + y`.
- A functional interface is an interface with a single abstract method that is used as the type of a lambda expression.

CHAPTER 4

Libraries

I've talked about how to write lambda expressions but so far haven't covered the other side of the fence: how to use them. This lesson is important even if you're not writing a heavily functional library like streams. Even the simplest application is still likely to have application code that could benefit from code as data.

Another Java 8 change that has altered the way that we need to think about libraries is the introduction of `default` methods and static methods on interfaces. This change means that methods on interfaces can now have bodies and contain code.

I'll also fill in some gaps in this chapter, covering topics such as what happens when you overload methods with lambda expressions and how to use primitives. These are important things to be aware of when you're writing lambda-enabled code.

Using Lambda Expressions in Code

In [Chapter 2](#), I described how a lambda expression is given the type of a functional interface and how this type is inferred. From the point of view of code calling the lambda expression, you can treat it identically to calling a method on an interface.

Let's look at a concrete example framed in terms of logging frameworks. Several commonly used Java logging frameworks, including `slf4j` and `log4j`, have methods that log output only when their logging level is set to a certain level or higher. So, they will have a method like `void debug(String message)` that will log `message` if the level is at `debug`.

Unfortunately, calculating the `message` to log frequently has a performance cost associated with it. Consequently, you end up with a situation in which people start explicitly calling the Boolean `isDebugEnabled` method in order to optimize this performance cost. A code sample is shown in [Example 4-1](#). Even though a direct call to `debug` would have avoided logging the text, it still would had to call the `expensiveOperation` method

and also concatenate its output to the message `String`, so the explicit `if` check still ends up being faster.

Example 4-1. A logger using `isDebugEnabled` to avoid performance overhead

```
Logger logger = new Logger();
if (logger.isDebugEnabled()) {
    logger.debug("Look at this: " + expensiveOperation());
}
```

What we actually want to be able to do is pass in a lambda expression that generates a `String` to be used as the message. This expression would be called only if the `Logger` was actually at debug level or above. This approach would allow us to rewrite the previous code example to look like the code in [Example 4-2](#).

Example 4-2. Using lambda expressions to simplify logging code

```
Logger logger = new Logger();
logger.debug(() -> "Look at this: " + expensiveOperation());
```

So how do we implement this method from within our `Logger` class? From the library point of view, we can just use the builtin `Supplier` functional interface, which has a single `get` method. We can then call `isDebugEnabled` in order to find out whether to call this method and pass the result into our `debug` method if it is enabled. The resulting code is shown in [Example 4-3](#).

Example 4-3. The implementation of a lambda-enabled logger

```
public void debug(Supplier<String> message) {
    if (isDebugEnabled()) {
        debug(message.get());
    }
}
```

Calling the `get()` method in this example corresponds to calling the lambda expression that was passed into the method to be called. This approach also conveniently works with anonymous inner classes, which allows you maintain a backward-compatible API if you have consumers of your code who can't upgrade to Java 8 yet.

It's important to remember that each of the different functional interfaces can have a different name for its actual method. So, if we were using a `Predicate`, we would have to call `test`, or if we were using `Function`, we would have to call `apply`.

Primitives

You might have noticed in the previous section that we skimmed over the use of *primitive* types. In Java we have a set of parallel types—for example, `int` and `Integer`—where one is a primitive type and the other a *boxed* type. Primitive types are built into the

language and runtime environment as fundamental building blocks; boxed types are just normal Java classes that wrap up the primitives.

Because Java generics are based around *erasing* a generic parameter—in other words, pretending it’s an instance of `Object`—only the boxed types can be used as generic arguments. This is why if you want a list of integer values in Java it will always be `List<Integer>` and not `List<int>`.

Unfortunately, because boxed types are objects, there is a memory overhead to them. For example, although an `int` takes 4 bytes of memory, an `Integer` takes 16 bytes. This gets even worse when you start to look at arrays of numbers, as each element of a primitive array is just the size of the primitive, while each element of a boxed array is actually an in-memory pointer to another object on the Java heap. In the worst case, this might make an `Integer[]` take up nearly six times more memory than an `int[]` of the same size.

There is also a computational overhead when converting from a primitive type to a boxed type, called *boxing*, and vice versa, called *unboxing*. For algorithms that perform lots of numerical operations, the cost of boxing and unboxing combined with the additional memory bandwidth used by allocated boxed objects can make the code significantly slower.

As a consequence of these performance overheads, the streams library differentiates between the primitive and boxed versions of some library functions. The `mapToLong` higher-order function and `ToLongFunction`, shown in [Figure 4-1](#), are examples of this effort. Only the `int`, `long`, and `double` types have been chosen as the focus of the primitive specialization implementation in Java 8 because the impact is most noticeable in numerical algorithms.




Figure 4-1. `ToLongFunction`

The primitive specializations have a very clear-cut naming convention. If the return type is a primitive, the interface is prefixed with `To` and the primitive type, as in `ToLongFunction` (shown in [Figure 4-1](#)). If the argument type is a primitive type, the name prefix is just the type name, as in `LongFunction` ([Figure 4-2](#)). If the higher-order function uses a primitive type, it is suffixed with `To` and the primitive type, as in `mapToLong`.




Figure 4-2. *LongFunction*

There are also specialized versions of `Stream` for these primitive types that prefix the type name, such as `LongStream`. In fact, methods like `mapToLong` don't return a `Stream`; they return these specialized streams. On the specialized streams, the `map` implementation is also specialized: it takes a function called `LongUnaryOperator`, visible in [Figure 4-3](#), which maps a `long` to a `long`. It's also possible to get back from a primitive stream to a boxed stream through higher-order function variations such as `mapToObj` and the `boxed` method, which returns a stream of boxed objects such as `Stream<Long>`.



Figure 4-3. *LongUnaryOperator*

It's a good idea to use the primitive specialized functions wherever possible because of the performance benefits. You also get additional functionality available on the specialized streams. This allows you to avoid having to implement common functionality and to use code that better conveys the intent of numerical operations. You can see an example of how to use this functionality in [Example 4-4](#).

Example 4-4. Using `summaryStatistics` to understand track length data

```

public static void printTrackLengthStatistics(Album album) {
    IntSummaryStatistics trackLengthStats
        = album.getTracks()
            .mapToInt(track -> track.getLength())
            .summaryStatistics();

    System.out.printf("Max: %d, Min: %d, Ave: %f, Sum: %d",
        trackLengthStats.getMax(),
        trackLengthStats.getMin(),
        trackLengthStats.getAverage(),
        trackLengthStats.getSum());
}
  
```

[Example 4-4](#) prints out a summary of track length information to the console. Instead of calculating that information ourselves, we map each track to its length, using the primitive specialized `mapToInt` method. Because this method returns an `IntStream`, we

can call `summaryStatistics`, which calculates statistics such as the minimum, maximum, average, and sum values on the `IntStream`.

These values are available on all the specialized streams, such as `DoubleStream` and `LongStream`. It's also possible to calculate the individual summary statistics if you don't need all of them through the `min`, `max`, `average`, and `sum` methods, which are all also available on all three primitive specialized `Stream` variants.

Overload Resolution

It's possible in Java to *overload* methods, so you have multiple methods with the same name but different signatures. This approach poses a problem for parameter-type inference because it means that there are several types that could be inferred. In these situations `javac` will pick the *most specific* type for you. For example, the method call in [Example 4-5](#), when choosing between the two methods in [Example 4-6](#), prints out `String`, not `Object`.

Example 4-5. A method that could be dispatched to one of two methods

```
overloadedMethod("abc");
```

Example 4-6. Two methods that are overloaded

```
private void overloadedMethod(Object o) {  
    System.out.print("Object");  
}  
  
private void overloadedMethod(String s) {  
    System.out.print("String");  
}
```

A `BinaryOperator` is special type of `BiFunction` for which the arguments and the return type are all the same. For example, adding two integers would be a `BinaryOperator`.

Because lambda expressions have the types of their functional interfaces, the same rules apply when passing them as arguments. We can overload a method with the `BinaryOperator` and an interface that extends it. When calling these methods, Java will infer the type of your lambda to be the most specific functional interface. For example, the code in [Example 4-7](#) prints out `IntegerBinaryOperator` when choosing between the two methods in [Example 4-8](#).

Example 4-7. Another overloaded method call

```
overloadedMethod((x, y) -> x + y);
```

Example 4-8. A choice between two overloaded methods

```
private interface IntegerBiFunction extends BinaryOperator<Integer> {  
  
}  
  
private void overloadedMethod(BinaryOperator<Integer> lambda) {  
    System.out.print("BinaryOperator");  
}  
  
private void overloadedMethod(IntegerBiFunction lambda) {  
    System.out.print("IntegerBinaryOperator");  
}
```

Of course, when there are multiple method overloads, there isn't always a clear "most specific type." Take a look at [Example 4-9](#).

Example 4-9. A compile failure due to overloaded methods

```
overloadedMethod((x) -> true);  
  
private interface IntPredicate {  
    public boolean test(int value);  
}  
  
private void overloadedMethod(Predicate<Integer> predicate) {  
    System.out.print("Predicate");  
}  
  
private void overloadedMethod(IntPredicate predicate) {  
    System.out.print("IntPredicate");  
}
```

The lambda expression passed into `overloadedMethod` is compatible with both a normal `Predicate` and the `IntPredicate`. There are method overloads for each of these options defined within this code block. In this case, `javac` will fail to compile the example, complaining that the lambda expression is an ambiguous method call: `IntPredicate` doesn't extend any `Predicate`, so the compiler isn't able to infer that it's more specific.

The way to fix these situations is to cast the lambda expression to either `IntPredicate` or `Predicate<Integer>`, depending upon which behavior you want to call. Of course, if you've designed the library yourself, you might conclude that this is a code smell and you should start renaming your overloaded methods.

In summary, the parameter types of a lambda are inferred from the *target type*, and the inference follows these rules:

- If there is a single possible target type, the lambda expression infers the type from the corresponding argument on the functional interface.
- If there are several possible target types, the most specific type is inferred.

- If there are several possible target types and there is no most specific type, you must manually provide a type.

@FunctionalInterface

Although I talked about the criteria for what a functional interface actually is back in [Chapter 2](#), I haven't yet mentioned the `@FunctionalInterface` annotation. This is an annotation that should be applied to any interface that is intended to be used as a functional interface.

What does that really mean? Well, there are some interfaces in Java that have only a single method but aren't normally meant to be implemented by lambda expressions. For example, they might assume that the object has internal state and be interfaces with a single method only coincidentally. A couple of good examples are `java.lang.Comparable` and `java.io.Closeable`.

If a class is `Comparable`, it means there is a defined order between instances, such as alphabetical order for strings. You don't normally think about functions themselves as being comparable objects because they lack fields and state, and if there are no fields and no state, what is there to sensibly compare?

For an object to be `Closeable` it must hold an open resource, such as a file handle that needs to be closed at some point in time. Again, the interface being called cannot be a pure function because closing a resource is really another example of mutating state.

In contrast to `Closeable` and `Comparable`, all the new interfaces introduced in order to provide `Stream` interoperability are expected to be implemented by lambda expressions. They are really there to bundle up blocks of code as data. Consequently, they have the `@FunctionalInterface` annotation applied.

Using the annotation compels `javac` to actually check whether the interface meets the criteria for being a functional interface. If the annotation is applied to an enum, class, or annotation, or if the type is an interface with more than one single abstract method, then `javac` will generate an error message. This is quite helpful for being able to catch errors easily when refactoring your code.

Binary Interface Compatibility

As you saw in [Chapter 3](#), one of the biggest API changes in Java 8 is to the collections library. As Java has evolved, it has maintained backward binary compatibility. In practical terms, this means that if you compiled a library or application with Java 1 through 7, it'll run out of the box in Java 8.

Of course, there are still bugs from time to time, but compared to many other programming platforms, binary compatibility has been viewed as a key Java strength. Barring the introduction of a new keyword, such as `enum`, there has also been an effort to maintain backward source compatibility. Here the guarantee is that if you’ve got source code in Java 1-7, it’ll compile in Java 8.

These guarantees are really hard to maintain when you’re changing such a core library component as the collections library. As a thought exercise, consider a concrete example. The `stream` method was added to the `Collection` interface in Java 8, which means that any class that implements `Collection` must also have this method on it. For core library classes, this problem can easily be solved by implementing that method (e.g., adding a `stream` method to `ArrayList`).

Unfortunately, this change still breaks binary compatibility because it means that any class outside of the JDK that implements `Collection`—say, `MyCustomList`—must also have implemented the `stream` method. In Java 8 `MyCustomList` would no longer compile, and even if you had a compiled version when you tried to load `MyCustomList` into a JVM, it would result in an exception being thrown by your `ClassLoader`.

This nightmare scenario of all third-party collections libraries being broken has been averted, but it did require the introduction of a new language concept: *default methods*.

Default Methods

So you’ve got your new `stream` method on `Collection`; how do you allow `MyCustomList` to compile without ever having to know about its existence? The Java 8 approach to solving the problem is to allow `Collection` to say, “If any of my children don’t have a `stream` method, they can use this one.” These methods on an interface are called *default methods*. They can be used on any interface, functional or not.

Another default method that has been added is the `forEach` method on `Iterable`, which provides similar functionality to the `for` loop but lets you use a lambda expression as the body of the loop. **Example 4-10** shows how this could be implemented in the JDK.

Example 4-10. An example default method, showing how `forEach` might be implemented

```
default void forEach(Consumer<? super T> action) {  
    for (T t : this) {  
        action.accept(t);  
    }  
}
```

Now that you’re familiar with the idea that you can use lambda expressions by just calling methods on interfaces, this example should look pretty simple. It uses a regular `for` loop to iterate over the underlying `Iterable`, calling the `accept` method with each value.

If it's so simple, why mention it? The important thing is that new `default` keyword right at the beginning of the code snippet. That tells `javac` that you really want to add a method to an interface. Other than the addition of a new keyword, default methods also have slightly different inheritance rules to regular methods.

The other big difference is that, unlike classes, interfaces don't have instance fields, so default methods can modify their child classes only by calling methods on them. This helps you avoid making assumptions about the implementation of their children.

Default Methods and Subclassing

There are some subtleties about the way that default methods override and can be overridden by other methods. Let's look the simplest case to begin with: no overriding. In [Example 4-11](#), our `Parent` interface defines a `welcome` method that sends a message when called. The `ParentImpl` class doesn't provide an implementation of `welcome`, so it inherits the default method.

Example 4-11. The `Parent` interface; the `welcome` method is a default

```
public interface Parent {  
  
    public void message(String body);  
  
    public default void welcome() {  
        message("Parent: Hi!");  
    }  
  
    public String getLastMessage();  
}
```

When we come to call this code, in [Example 4-12](#), the default method is called and our assertion passes.

Example 4-12. Using the default method from client code

```
@Test  
public void parentDefaultUsed() {  
    Parent parent = new ParentImpl();  
    parent.welcome();  
    assertEquals("Parent: Hi!", parent.getLastMessage());  
}
```

Now we can extend `Parent` with a `Child` interface, whose code is listed in [Example 4-13](#). `Child` implements its own default `welcome` method. As you would intuitively expect, the default method on `Child` overrides the default method on `Parent`. In this example, again, the `ChildImpl` class doesn't provide an implementation of `welcome`, so it inherits the default method.

Example 4-13. Child interface that extends Parent

```
public interface Child extends Parent {  
  
    @Override  
    public default void welcome() {  
        message("Child: Hi!");  
    }  
  
}
```

You can see the class hierarchy at this point in Figure 4-4.

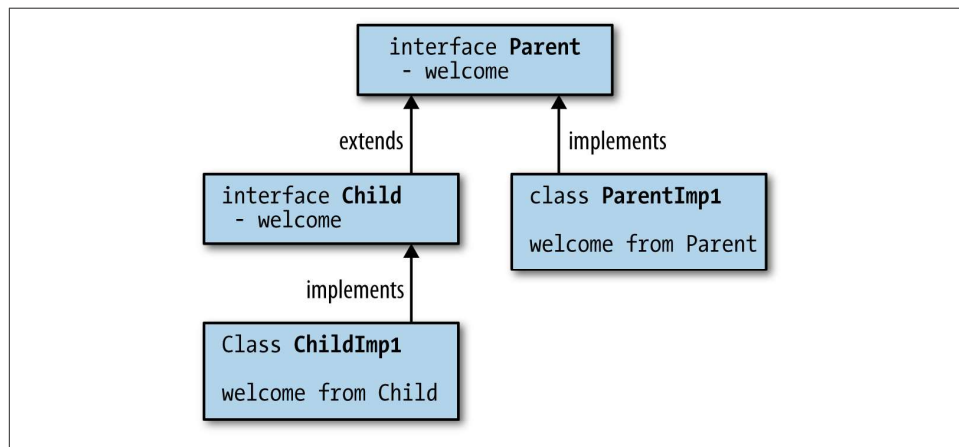


Figure 4-4. A diagram showing the inheritance hierarchy at this point

Example 4-14 calls this interface and consequently ends up sending the string "Child: Hi!".

Example 4-14. Client code that calls our Child interface

```
@Test  
public void childOverrideDefault() {  
    Child child = new ChildImpl();  
    child.welcome();  
    assertEquals("Child: Hi!", child.getLastMessage());  
}
```

Now the default method is a virtual method—that is, the opposite of a static method. What this means is that whenever it comes up against competition from a class method, the logic for determining which override to pick always chooses the class. A simple example of this is shown in Examples 4-15 and 4-16, where the welcome method of OverridingParent is chosen over that of Parent.

Example 4-15. A parent class that overrides the default implementation of welcome

```
public class OverridingParent extends ParentImpl {  
  
    @Override  
    public void welcome() {  
        message("Class Parent: Hi!");  
    }  
  
}
```

Example 4-16. An example of a concrete method beating a default method

```
@Test  
public void concreteBeatsDefault() {  
    Parent parent = new OverridingParent();  
    parent.welcome();  
    assertEquals("Class Parent: Hi!", parent.getLastMessage());  
}
```

Here's a situation, presented in [Example 4-18](#), in which you might not expect the concrete class to override the default method. `OverridingChild` inherits both the `welcome` method from `Child` and the `welcome` method from `OverridingParent` and doesn't do anything itself. `OverridingParent` is chosen despite `OverridingChild` (the code in [Example 4-17](#)), being a more specific type because it's a concrete method from a class rather than a default method (see [Figure 4-5](#)).

Example 4-17. Again, our child interface overrides the default welcome method

```
public class OverridingChild extends OverridingParent implements Child {  
  
}
```

Example 4-18. An example of a concrete method beating a default method that is more specific

```
@Test  
public void concreteBeatsCloserDefault() {  
    Child child = new OverridingChild();  
    child.welcome();  
    assertEquals("Class Parent: Hi!", child.getLastMessage());  
}
```

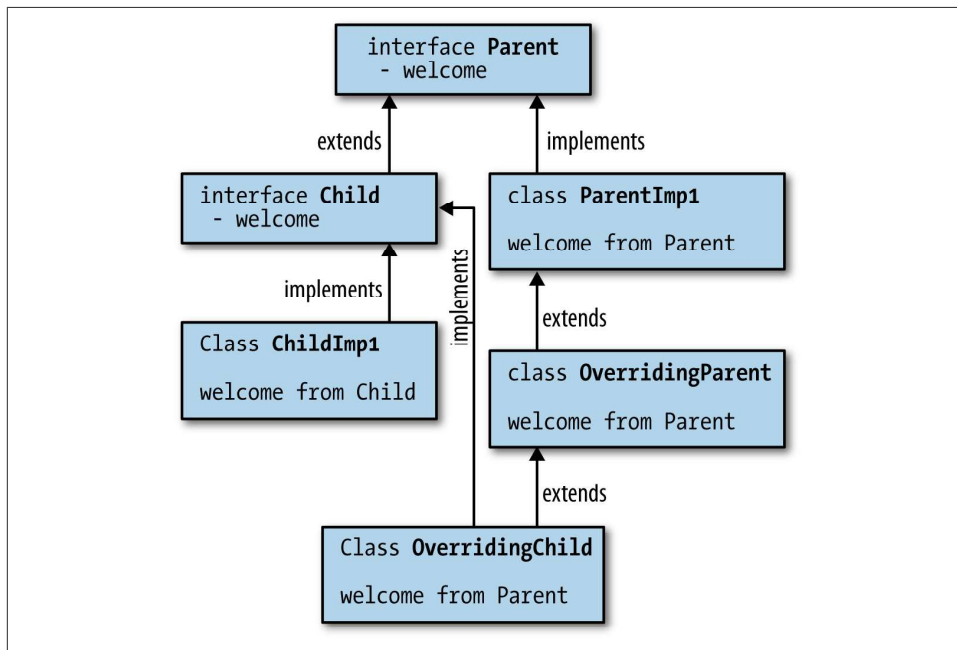


Figure 4-5. A diagram showing the complete inheritance hierarchy

Put simply: *class wins*. The motivation for this decision is that default methods are designed primarily to allow binary compatible API evolution. Allowing classes to win over any default methods simplifies a lot of inheritance scenarios.

Suppose we had a custom list implementation called `MyCustomList` and had implemented a custom `addAll` method, and the new `List` interface provided a default `addAll` that delegated to the `add` method. If the default method wasn't guaranteed to be overridden by this `addAll` method, we could break the existing implementation.

Multiple Inheritance

Because interfaces are subject to multiple inheritance, it's possible to get into situations where two interfaces both provide default methods with the same signature. Here's an example in which both a `Carriage` and a `Jukebox` provide a method to `rock`—in each case, for different purposes. We also have a `MusicalCarriage`, which is both a `Jukebox` (Example 4-19) and a `Carriage` (Example 4-20) and tries to inherit the `rock` method.

Example 4-19. *Jukebox*

```

public interface Jukebox {

    public default String rock() {

```

```

        return "... all over the world!";
    }
}

```

Example 4-20. Carriage

```

public interface Carriage {

    public default String rock() {
        return "... from side to side";
    }

}

public class MusicalCarriage implements Carriage, Jukebox {
}

```

Because it's not clear to javac which method it should inherit, this will just result in the compile error class MusicalCarriage inherits unrelated defaults for rock() from types Carriage and Jukebox. Of course, it's possible to resolve this by implementing the rock method, shown in [Example 4-21](#).

Example 4-21. Implementing the rock method

```

public class MusicalCarriage
    implements Carriage, Jukebox {

    @Override
    public String rock() {
        return Carriage.super.rock();
    }

}

```

This example uses the enhanced super syntax in order to pick Carriage as its preferred rock implementation. Previously, super acted as a reference to the parent class, but by using the InterfaceName.super variant it's possible to specify a method from an inherited interface.

The Three Rules

If you're ever unsure of what will happen with default methods or with multiple inheritance of behavior, there are three simple rules for handling conflicts:

1. Any class wins over any interface. So if there's a method with a body, or an abstract declaration, in the superclass chain, we can ignore the interfaces completely.

2. Subtype wins over supertype. If we have a situation in which two interfaces are competing to provide a `default` method and one interface extends the other, the subclass wins.
3. *No rule 3.* If the previous two rules don't give us the answer, the subclass must either implement the method or declare it `abstract`.

Rule 1 is what brings us compatibility with old code.

Tradeoffs

These changes raise a bunch of issues regarding what an interface really is in Java 8, as you can define methods with code bodies on them. This means that interfaces now provide a form of multiple inheritance that has previously been frowned upon and whose removal has been considered a usability advantage of Java over C++.

No language feature is always good or always bad. Many would argue that the real issue is multiple inheritance of state rather than just blocks of code, and as `default` methods avoid multiple inheritance of state, they avoid the worst pitfalls of multiple inheritance in C++.

It can also be very tempting to try and work around these limitations. Blog posts have already cropped up trying to implement full-on traits with multiple inheritance of state as well as `default` methods. Trying to hack around the deliberate restrictions of Java 8 puts us back into the old pitfalls of C++.

It's also pretty clear that there's still a distinction between interfaces and abstract classes. Interfaces give you multiple inheritance but no fields, while abstract classes let you inherit fields but you don't get multiple inheritance. When modeling your problem domain, you need to think about this tradeoff, which wasn't necessary in previous versions of Java.

Static Methods on Interfaces

We've seen a lot of calling of `Stream.of` but haven't gotten into its details yet. You may recall that `Stream` is an interface, but this is a static method on an interface. This is another new language change that has made its way into Java 8, primarily in order to help library developers, but with benefits for day-to-day application developers as well.

An idiom that has accidentally developed over time is ending up with classes full of static methods. Sometimes a class can be an appropriate location for utility code, such as the `Objects` class introduced in Java 7 that contained functionality that wasn't specific to any particular class.

Of course, when there's a good semantic reason for a method to relate to a concept, it should always be put in the same class or interface rather than hidden in a utility class to the side. This helps structure your code in a way that's easier for someone reading it to find the relevant method.

For example, if you want to create a simple `Stream` of values, you would expect the method to be located on `Stream`. Previously, this was impossible, and the addition of a very interface-heavy API in terms of `Stream` finally motivated the addition of static methods on interfaces.



There are other methods on `Stream` and its primitive specialized variants. Specifically, `range` and `iterate` give us other ways of generating our own streams.

Optional

Something I've glossed over so far is that `reduce` can come in a couple of forms: the one we've seen, which takes an initial value, and another variant, which doesn't. When the initial value is left out, the first call to the reducer uses the first two elements of the `Stream`. This is useful if there's no sensible initial value for a `reduce` operation and will return an instance of `Optional`.

`Optional` is a new core library data type that is designed to provide a better alternative to `null`. There's quite a lot of hatred for the old `null` value. Even the man who invented the concept, Tony Hoare, described it as “my billion-dollar mistake.” That's the trouble with being an influential computer scientist—you can make a billion-dollar mistake without even seeing the billion dollars yourself!

`null` is often used to represent the absence of a value, and this is the use case that `Optional` is replacing. The problem with using `null` in order to represent absence is the dreaded `NullPointerException`. If you refer to a variable that is `null`, your code blows up. The goal of `Optional` is twofold. First, it encourages the coder to make appropriate checks as to whether a variable is `null` in order to avoid bugs. Second, it documents values that are expected to be absent in a class's API. This makes it easier to see where the bodies are buried.

Let's take a look at the API for `Optional` in order to get a feel for how to use it. If you want to create an `Optional` instance from a value, there is a factory method called `of`. The `Optional` is now a container for this value, which can be pulled out with `get`, as shown in [Example 4-22](#).

Example 4-22. Creating an Optional from a value

```
Optional<String> a = Optional.of("a");
assertEquals("a", a.get());
```

Because an `Optional` may also represent an absent value, there's also a factory method called `empty`, and you can convert a nullable value into an `Optional` using the `ofNullable` method. Both of these are shown in [Example 4-23](#), along with the use of the `isPresent` method (which indicates whether the `Optional` is holding a value).

Example 4-23. Creating an empty Optional and checking whether it contains a value

```
Optional emptyOptional = Optional.empty();
Optional alsoEmpty = Optional.ofNullable(null);

assertFalse(emptyOptional.isPresent());

// a is defined above
assertTrue(a.isPresent());
```

One approach to using `Optional` is to guard any call to `get()` by checking `isPresent()`. A neater approach is to call the `orElse` method, which provides an alternative value in case the `Optional` is empty. If creating an alternative value is computationally expensive, the `orElseGet` method should be used. This allows you to pass in a `Supplier` that is called only if the `Optional` is genuinely empty. Both of these methods are demonstrated in [Example 4-24](#).

Example 4-24. Using orElse and orElseGet

```
assertEquals("b", emptyOptional.orElse("b"));
assertEquals("c", emptyOptional.orElseGet(() -> "c"));
```

Not only is `Optional` used in new Java 8 APIs, but it's also just a regular class that you can use yourself when writing domain classes. This is definitely something to think about when trying to avoid nullness-related bugs such as uncaught exceptions.

Key Points

- A significant performance advantage can be had by using primitive specialized lambda expressions and streams such as `IntStream`.
- Default methods are methods with bodies on interfaces prefixed with the keyword `default`.
- The `Optional` class lets you avoid using `null` by modeling situations where a value may not be present.