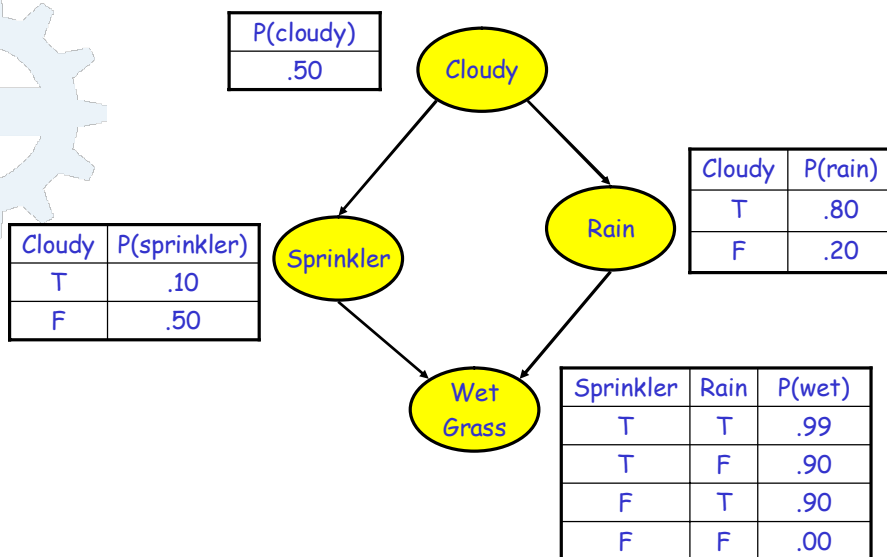


## 14.5 Approximate Inference

- Given the intractability of exact inference, it is essential to consider approximate inference methods
- Approximation is based on random sampling from a known probability distribution (Monte Carlo algorithms)
- E.g., an unbiased coin can be thought of as a random variable *Coin* with values [heads, tails] and a prior distribution  
 $P(\text{Coin}) = [0.5, 0.5]$
- Sampling from this distribution is exactly like flipping the coin: with probability 0.5 it will return heads, and with probability 0.5 it will return tails
- Given a source of random numbers in the range [0, 1], it is a simple matter to sample any distribution on a single variable



### 14.5.1 Direct Sampling Methods

- From a Bayesian network that has no evidence associated with it, we can sample each variable in turn, in topological order
- When the values of parent nodes have been drawn, we know from which distribution we have to sample the child
- Let us fix a topological order for the nodes of our network:  
[Cloudy, Sprinkler, Rain, WetGrass]
  1. Sample from  $P(\text{Cloudy}) = [0.5, 0.5]$ ; suppose this returns **True**
  2. Sample from  $P(\text{Sprinkler} \mid \text{cloudy}) = [0.1, 0.9]$ ; suppose this returns **False**
  3. Sample from  $P(\text{Rain} \mid \text{cloudy}) = [0.8, 0.2]$ ; suppose this returns **True**
  4. Sample from  $P(\text{WetGrass} \mid \neg \text{sprinkler}, \text{rain}) = [0.9, 0.1]$ ; suppose this returns **True**



- From the prior joint distribution specified by the network we have drawn the event [True, False, True, True]
- Let  $S_{PS}(x_1, \dots, x_n)$  be the probability that a specific event is generated by this prior sampling algorithm
- Just looking at the sampling process, we have
 
$$S_{PS}(x_1, \dots, x_n) = \prod_{i=1, \dots, n} P(x_i \mid \text{parents}(X_i))$$
- On the other hand, this is also the probability of the event according to the Bayesian net's representation of the joint distribution; i.e.:
 
$$S_{PS}(x_1, \dots, x_n) = P(x_1, \dots, x_n)$$
- Let  $N_{PS}(x_1, \dots, x_n)$  be the frequency of the specific event  $x_1, \dots, x_n$  and that there are  $N$  total samples



- We expect this frequency to converge in the limit to its expected value according to the sampling probability:

$$\lim_{N \rightarrow \infty} N_{PS}(x_1, \dots, x_n)/N = S_{PS}(x_1, \dots, x_n) = P(x_1, \dots, x_n)$$

- E.g.,  $S_{PS}([True, False, True, True]) = 0.5 \times 0.9 \times 0.8 \times 0.9 = 0.324$ , hence in the limit of large  $N$ , we expect 32.4% of the samples to be of this event
- The estimate of prior sampling is *consistent* in the sense that the estimated probability becomes exact in the large-sample limit
- One can also produce a consistent estimate of the probability of any partially specified event  $x_1, \dots, x_m$ , where  $m \leq n$ :  

$$P(x_1, \dots, x_m) \approx N_{PS}(x_1, \dots, x_m)/N$$
- Let us denote by  $P_{data}(\cdot)$  the probability estimated from a sample



## Rejection Sampling

- To determine a conditional probability  $P(X | e)$  we could apply the following simple sampling approach:
  1. Generate samples from the prior distribution specified by the network
  2. Reject all those that do not match the evidence  $e$
  3. The estimate  $P_{data}(X = x | e)$  is obtained by counting how often  $X = x$  occurs in the remaining samples
- The estimated distribution  $P_{data}(X | e)$  that the algorithm returns is, by the definition of the algorithm  

$$\propto N_{PS}(X, e) = N_{PS}(X, e) / N_{PS}(e)$$
- As an estimate of the probability of a partially specified event it is consistent

$$P_{data}(X | e) \approx P(X, e) / P(e) = P(X | e)$$



- Let us generate 100 samples in order to estimate the distribution  $P(\text{Rain} \mid \text{Sprinkler} = \text{True})$ 
  - Suppose that 73 of those that we generate have  $\text{Sprinkler} = \text{False}$  and are rejected
  - The remaining 27 have  $\text{Sprinkler} = \text{True}$
  - Out of them 8 have  $\text{Rain} = \text{True}$  and 19 have  $\text{Rain} = \text{False}$
- Hence, we now have  $P_{\text{data}}(\text{Rain} \mid \text{sprinkler}) \approx [0.296, 0.704]$ , while the true distribution is  $[0.3, 0.7]$
- As more samples are collected, the estimate will converge to the true answer
- The standard deviation of the error in each probability will be proportional to  $1/\sqrt{n}$ , where  $n$  is the number of samples
- The large number of rejected samples is a big problem:  
The fraction of samples consistent with the evidence drops exponentially as the number of evidence variables grows



## Likelihood weighting

- Rejection sampling is inefficient because it ends up rejecting so many of the generated samples
- To avoid generating needless samples that anyhow get rejected, let us fix the values for the evidence variables  $E$  and sample only the remaining variables  $X$  and  $Y$
- Not all events are equal, however
- Each event is weighted by the *likelihood* that the event accords to the evidence
- The likelihood is measured by the product of the conditional probabilities for each evidence variable, given its parents
- Intuitively, events in which the actual evidence appears unlikely should be given less weight





- To answer the query  $P(\text{Rain} \mid \text{sprinkler}, \text{wetgrass})$ , the weight  $w$  is first set to 1.0
- Sample from  $P(\text{Cloudy}) = [0.5, 0.5]$ ; suppose this returns **True**
- **Sprinkler** is an evidence variable with value **True**, therefore we update the weight
 
$$w \leftarrow w \times P(\text{sprinkler} \mid \text{cloudy}) = 0.1$$
- Sample from  $P(\text{Rain} \mid \text{cloudy}) = [0.8, 0.2]$ ; suppose this returns **True**
- **WetGrass** is an evidence variable with value **True**  $\Rightarrow$ 

$$w \leftarrow w \times P(\text{wetgrass} \mid \text{sprinkler}, \text{rain}) = 0.099$$
- Hence, the algorithm returns the event **[True, True, True, True]** with weight **0.099** and this is tallied under **Rain = True**



- Let us denote  $\mathbf{Z} = \{ \mathbf{X} \} \cup \mathbf{Y}$
- The weighted sample algorithm samples each variable in  $\mathbf{Z}$  given its parent values
 
$$S_{WS}(\mathbf{z}, \mathbf{e}) = \prod_{i=1, \dots, I} P(z_i \mid \text{parents}(Z_i))$$
- $\text{Parents}(Z_i)$  can include both hidden variables and evidence variables
- The sampling distribution  $S_{WS}$  pays some attention to the evidence, unlike the prior distribution  $P(\mathbf{z})$
- In  $S_{WS}$  the sampled values for each  $Z_i$  will be influenced by evidence among  $Z_i$ 's ancestors
- On the other hand, the true posterior distribution  $P(\mathbf{z} \mid \mathbf{e})$  also takes non-ancestor evidence into account



- The likelihood weight  $w$  makes up for the difference between the actual and desired sampling distributions

- Let a given sample  $\mathbf{x}$  be composed from  $\mathbf{z}$  and  $\mathbf{e}$ , then

$$w(\mathbf{z}, \mathbf{e}) = \prod_{i=1, \dots, m} P(e_i \mid \text{parents}(E_i))$$

- The *weighted* probability of a sample,  $S_{WS}(\mathbf{z}, \mathbf{e}) \cdot w(\mathbf{z}, \mathbf{e})$ , is

$$\prod_{i=1, \dots, l} P(z_i \mid \text{parents}(Z_i)) \cdot \prod_{i=1, \dots, m} P(e_i \mid \text{parents}(E_i)) \\ = P(\mathbf{z}, \mathbf{e}),$$

because the two products cover all the variables in the network



- Now it is easy to show that likelihood weighting estimates are consistent:

$$\begin{aligned} P_{\text{data}}(\mathbf{x} \mid \mathbf{e}) &= \alpha \sum_{\mathbf{y}} N_{WS}(\mathbf{x}, \mathbf{y}, \mathbf{e}) w(\mathbf{x}, \mathbf{y}, \mathbf{e}) && \text{algorithm} \\ &\approx \alpha' \sum_{\mathbf{y}} S_{WS}(\mathbf{x}, \mathbf{y}, \mathbf{e}) w(\mathbf{x}, \mathbf{y}, \mathbf{e}) && \text{for large } N \\ &= \alpha' \sum_{\mathbf{y}} P(\mathbf{x}, \mathbf{y}, \mathbf{e}) && \text{by prev. slide} \\ &= \alpha' P(\mathbf{x}, \mathbf{e}) \\ &= P(\mathbf{x} \mid \mathbf{e}) \end{aligned}$$

- Because likelihood weighting uses all the samples generated, it can be much more efficient than rejection sampling
- It will, however, suffer a degradation in performance as the number of evidence variables increases
- Because most samples will have very low weights, the weighted estimate will be dominated by a tiny fraction of samples



## 14.5.2 Inference by Markov Chain Simulation

- Markov chain Monte Carlo (MCMC)
- A Monte Carlo algorithm is a randomized algorithm, which can give the false answer with a small probability (vs. Las Vegas algorithm)
- MCMC generates each event by making a random change to the preceding event
- The next state is generated by randomly sampling a value for one of the nonevidence variables  $X_i$ , conditioned on the current values in its *Markov blanket*
- The Markov blanket of a variable consists of its parents, children, and children's parents
- MCMC therefore wanders randomly around the state space flipping one variable at a time, but keeping the evidence variables fixed



## 16 MAKING SIMPLE DECISIONS

- Let us associate each state  $S$  with a numeric *utility*  $U(S)$ , which expresses the desirability of the state
- A nondeterministic action  $a$  will have possible outcome states  $\text{Result}(a) = s'$
- Prior to the execution of  $a$  the agent assigns probability  $P(\text{Result}(a) = s' \mid a, e)$  to each outcome, where  $e$  summarizes the agent's available evidence of the world
- The expected utility of  $a$  can now be calculated:

$$EU(a \mid e) = \sum_{s'} P(\text{Result}(a) = s' \mid a, e) \cdot U(s')$$



- The principle of *maximum expected utility* (MEU) says that a rational agent should choose an action that maximizes the agent's expected utility  $\text{argmax}_a EU(a | e)$
- If we wanted to choose the best sequence of actions using this equation, we would have to enumerate all action sequences, which is clearly infeasible for long sequences
- If the utility function correctly reflects the performance measure by which the behavior is being judged  $\Rightarrow$  using MEU the agent will achieve the highest possible performance score averaged over the environments in which it could be placed
- Let us model a nondeterministic action with a *lottery*  $L$ , where possible outcomes  $S_1, \dots, S_n$  can occur with probabilities  $p_1, \dots, p_n$   

$$L = [p_1, S_1; p_2, S_2; \dots; p_n, S_n]$$



## 16.2 The Basis of Utility Theory

- $A \succ B$  Agent prefers lottery  $A$  over  $B$
- $A \sim B$  The agent is indifferent between  $A$  and  $B$
- $A \succeq B$  The agent prefers  $A$  to  $B$  or is indifferent between them

- Deterministic lottery  $[1, A] \equiv A$
- Reasonable constraints on the preference relation (in the name of rationality)
  - **Orderability:** given any two states, a rational agent must either prefer one to the other or else rate the two as equally preferable.

$$(A \succ B) \vee (B \succ A) \vee (A \sim B)$$

- **Transitivity:**

$$(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$$





- **Continuity:**

$$A \succ B \succ C \Rightarrow \exists p: [p, A; 1-p, C] \sim B$$

- **Substitutability:**

$$A \sim B \Rightarrow [p, A; 1-p, C] \sim [p, B; 1-p, C]$$

- **Monotonicity:**

$$A \succ B \Rightarrow (p \geq q \Leftrightarrow [p, A; 1-p, B] \succeq [q, A; 1-q, B])$$

- **Decomposability:** Compound lotteries can be reduced to simpler ones by the laws of probability

$$[p, A; 1-p, [q, B; 1-q, C]] \sim [p, A; (1-p)q, B; (1-p)(1-q), C]$$

- Notice that these axioms of utility theory do not say anything about utility
- The existence of a utility function follows from them



## Preferences lead to utility

### 1. Existence of Utility Function:

If an agent's preferences follow the axioms of utility, then there exists a real-valued function  $U$  s.t.

$$U(A) > U(B) \Leftrightarrow A \succ B$$

$$U(A) = U(B) \Leftrightarrow A \sim B$$

### 2. Expected Utility of a Lottery:

The utility of a lottery is

$$U([p_1, S_1; \dots; p_n, S_n]) = \sum_{i=1, \dots, n} p_i \cdot U(S_i)$$

Because the outcome of a nondeterministic action is a lottery, this gives us the MEU decision rule from slide 254





- The axioms of utility do not specify a unique utility function for an agent
- For example, we can transform a utility function  $U(S)$  into
 
$$U'(S) = aU(S) + b$$
 where  $b$  is a constant and  $a$  is any positive constant
- Clearly, this affine transformation leaves the agent's behavior unchanged
- In deterministic contexts, where there are states but no lotteries, behavior is unchanged by any monotonic transformation
- E.g., the cube root of the utility  $\sqrt[3]{U(S)}$
- Utility function is ordinal — it really provides just rankings of states rather than meaningful numerical values



### 16.3 Utility Functions

- Money (or an agent's total net assets) would appear to be a straightforward utility measure
- The agent exhibits a monotonic preference for definite amounts of money
- We need to determine a model for lotteries involving money
  - We have won a million euros in a TV game show
  - The host offers to flip a coin, if the coin comes up heads, we end up with nothing, but if it comes up tails, we win three million euros
  - Is the only rational choice to accept the offer which has the expected monetary value of 1,5 million euros?
- The true question is maximizing total wealth (not winnings)



## Normalized utilities

- The scale of utilities reaches from the best possible prize  $u_T$  to the worst possible catastrophe  $u_L$
- *Normalized utilities* use a scale with  $u_L = 0$  and  $u_T = 1$
- Utilities of intermediate outcomes are assessed by asking the agent to indicate a preference between the given outcome state  $S$  and a standard lottery  $[p, u_T; 1-p, u_L]$
- The probability  $p$  is adjusted until the agent is indifferent between  $S$  and the standard lottery
- Assuming normalized utilities, the utility of  $S$  is given by  $p$




## 16.4 Multiattribute Utility Functions

- Most often the utility is determined by the values  $\mathbf{x} = [x_1, \dots, x_n]$  of multiple variables (attributes)  $\mathbf{X} = X_1, \dots, X_n$
- For simplicity, we will assume that each attribute is defined in such a way that, all other things being equal, higher values of the attribute correspond to higher utilities
- If for a pair of attribute vectors  $\mathbf{x}$  and  $\mathbf{y}$  it holds that  $x_i \geq y_i \quad \forall i$ , then  $\mathbf{x}$  *strictly dominates*  $\mathbf{y}$
- Suppose that airport site  $S_1$  costs less, generates less noise pollution, and is safer than site  $S_2$ , one would not hesitate to reject the latter
- In the general case, where the action outcomes are uncertain, strict dominance occurs less often than in the deterministic case
- *Stochastic dominance* is more useful generalization



- Suppose we believe that the cost of siting an airport is uniformly distributed between
  - $S_1$ : 2.8 and 4.8 billion euros
  - $S_2$ : 3.0 and 5.2 billion euros
- Then by examining the cumulative distributions, we see that  $S_1$  stochastically dominates  $S_2$  (because costs are negative)



- Cumulative distribution integrates the original distribution
- If two actions  $A_1$  and  $A_2$  lead to probability distributions  $p_1(x)$  and  $p_2(x)$  on attribute  $X$
- $A_1$  stochastically dominates  $A_2$  on  $X$  if
 
$$\forall x: \int_{-\infty, \dots, x} p_1(x') dx' \leq \int_{-\infty, \dots, x} p_2(x') dx'$$
- If
  - $A_1$  stochastically dominates  $A_2$ ,
  - then for any monotonically nondecreasing utility function  $U(x)$ ,  
the expected utility of  $A_1$  is at least as high as that of  $A_2$
- Hence, if an action is stochastically dominated by another action on all attributes, then it can be discarded



## 16.6 The Value of Information

- BP is hoping to buy one of  $n$  indistinguishable blocks of ocean drilling rights at the Gulf of Mexico
- Exactly one of the blocks contains oil worth  $C$  euros
- The price for each block is  $C/n$  euros
- A seismologist offers BP the results of a survey of block #3, which indicates definitively whether the block contains oil
- How much should BP be willing to pay for the information?
  - With probability  $1/n$ , the survey will indicate oil in block #3, in which case BP will buy the block for  $C/n$  euros and make a profit of  $(n-1)C/n$  euros
  - With probability  $(n-1)/n$ , the survey will show that the block contains no oil, in which case BP will buy a different block



- Now the probability of finding oil in one of the other blocks changes to  $1/(n-1)$ , so BP makes an expected profit of  $C/(n-1) - C/n = C/n(n-1)$  euros
- Now we can calculate the expected profit, given the survey information:
 
$$(1/n) \cdot ((n-1)C/n) + ((n-1)/n) \cdot (C/n(n-1)) = C/n$$
- Therefore, BP should be willing to pay the seismologist up to the price of the block itself
- With the information, one's course of action can be changed to suit the actual situation
- Without the information, one has to do what's best on average over the possible situations

